Two really cool concepts at the boundary of RL

Part I - Meta RL

Part II: safe shallow model-based RL

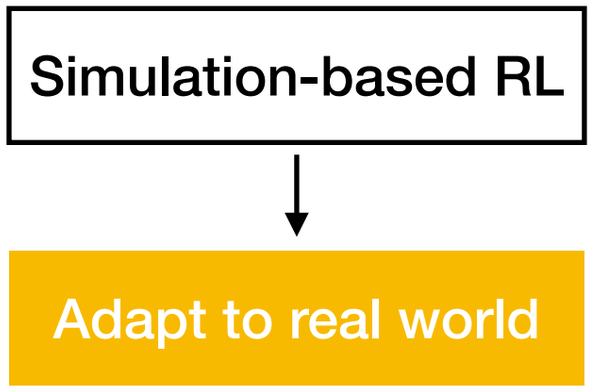

- Meta RL
- Adapts quickly to changes
- Brings nice properties

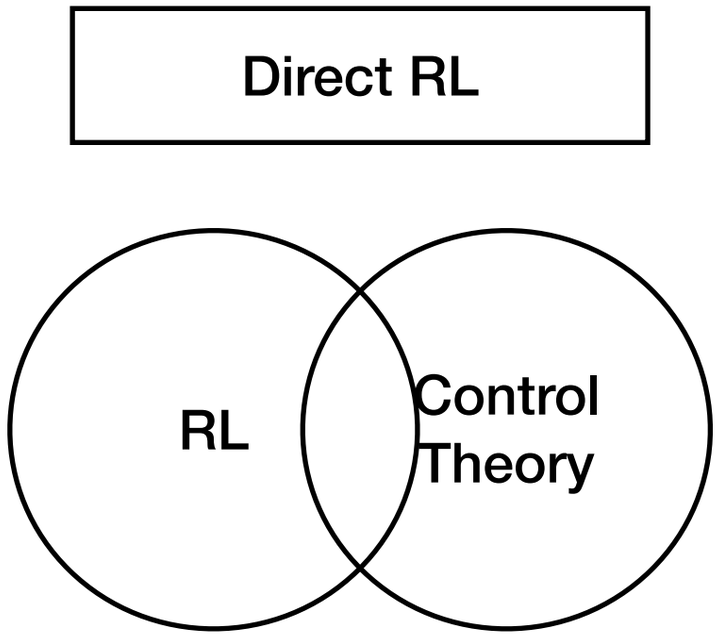

RL towards control theory - the BO of RL

Optimisation

- Optimisation has become a standard tool in the control room:
  - ➡ Fast adaption from scratch
  - ➡ Easy to tune with short exploration
  - ➡ It is not RL - optimisation is greedy
- RL has potential to solve a much broader range of problems:
  - ➡ Incorporates state information - if trained, much faster than optimization
  - ➡ Can handle delayed consequences
  - ➡ Policy might be faster and easier to calculate and implement