

RL and optimization

- All said falls into the domain of optimization:
 - ➔ An optimiser tries to find the arguments of a function to maximise the function value (optimization is greedy!)
 - ➔ RL algorithms look to find a mapping (the policy) from states to actions maximising the expected cumulative reward rather than just a single optimal function value
 - If parametric function approximation is used, we try to find the values of the parameters of the approximated function (either a value function or the policy directly) to obtain this mapping (this a classical optimisation problem).
- RL is comparable to calculus of variation (its origin is in classical mechanics - HJB equation) instead of function optimization

Optimization

$$\begin{aligned} &\text{maximise}_{\{A_i\}} \sum_{t=0}^T R(S_t, A_t, W_t) \\ &\text{subject to: } S_{t+1} = f(S_t, A_t, W_t) \end{aligned}$$

RL

$$\begin{aligned} &\text{maximise}_{\pi_t} \mathbb{E}_{W_t} \left[\sum_{t=0}^T R_t(S_t, A_t, W_t) \right] \\ &\text{subject to: } S_{t+1} = f(S_t, A_t, W_t) \\ &\quad A_t = \pi_t(S_t, S_{t-1}, \dots) \end{aligned}$$

Feedback structure takes noise into account

Often optimization is performed only for one step horizon:

$$\text{maximise}_a R(\cdot, a, W_t)$$

Possible solutions

- Use a high fidelity model
 - ➔ Usually unavailable...
- Development of faster algorithms
 - ➔ SAC (very efficient maximum entropy algorithm)
 - ➔ TD3 (simple and efficient tricks accelerate DDPG-Style algorithms)
- Use data efficient method as MBRL
 - ➔ Hard to handle, low computational efficiency
- Reuse von prior knowledge to accelerate RL
 - ➔ RL2: Fast reinforcement learning via slow reinforcement learning <https://arxiv.org/pdf/1611.02779.pdf>
 - ➔ Learning to reinforcement learning <https://arxiv.org/pdf/1611.05763.pdf>
 - ➔ Model-agnostic meta-learning <http://proceedings.mlr.press/v70/finn17a/finn17a.pdf>