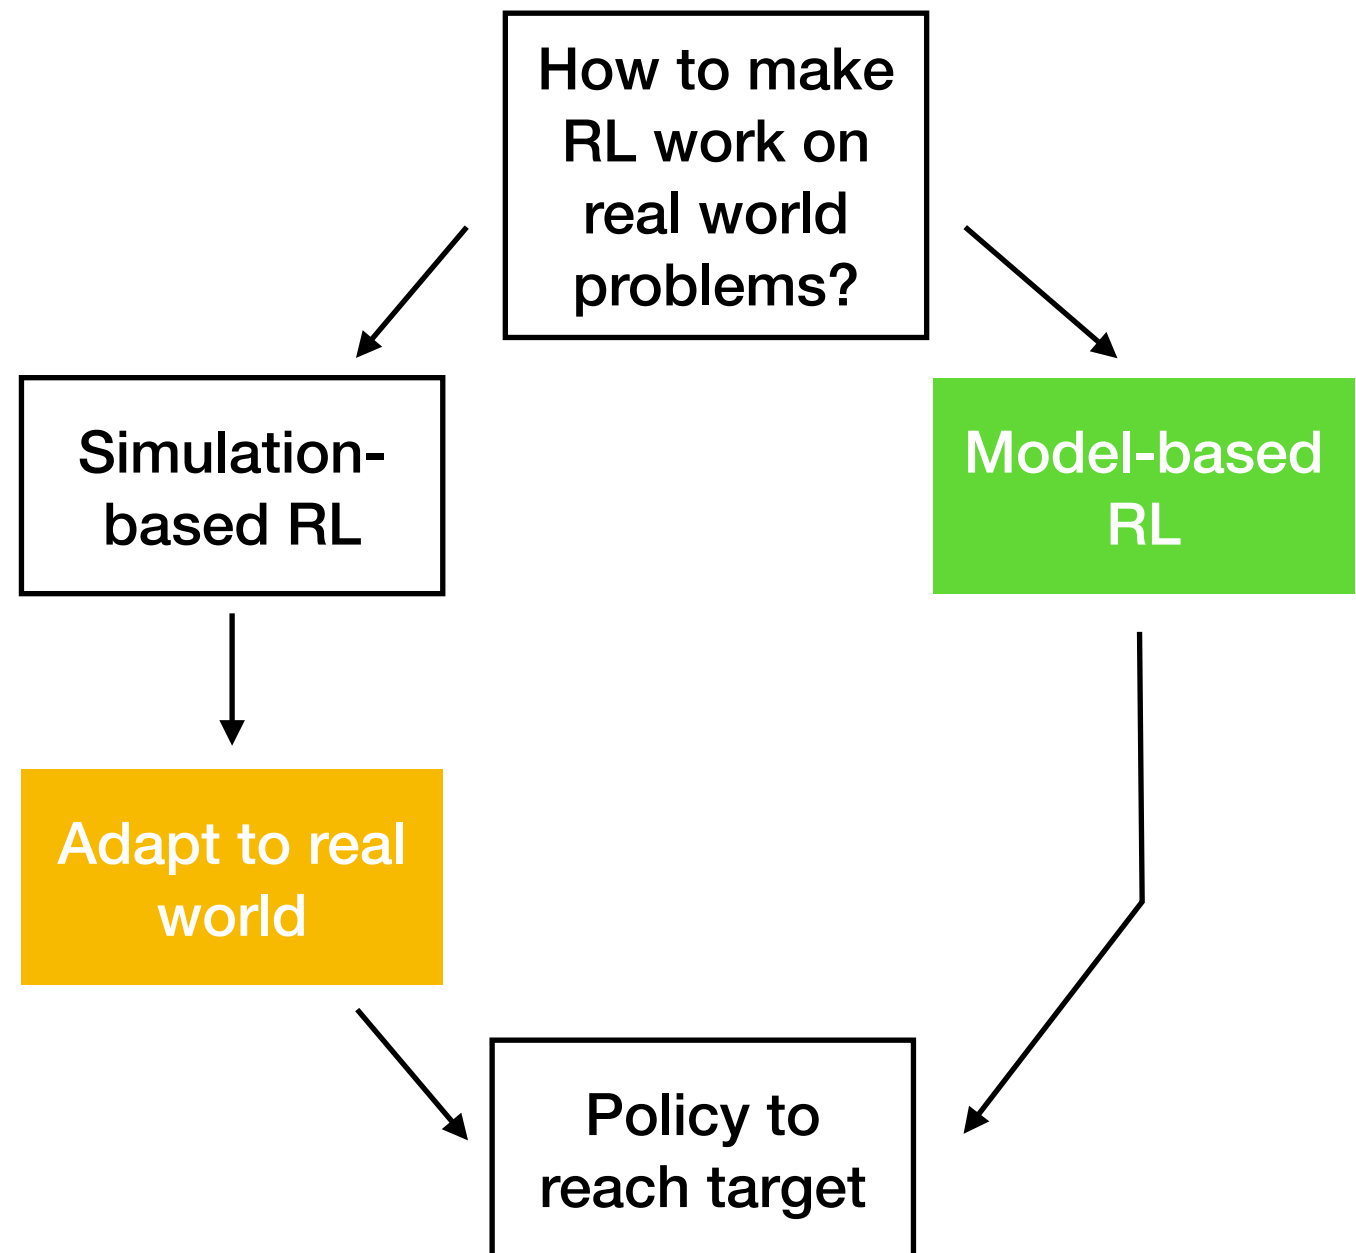


Scenarios

- If we have a simulation: prepare agent for real world training in the best way
- If we don't have a simulation: Train the agent in the fastest best way
- If we trained an agent...
 - ➔ Works on one day but not on others...
- Agent should be able to adapt quickly



Two really cool concepts at the boundary of RL

Part I - Meta RL

Simulation-based RL




Adapt to real world

- Meta RL
- Adapts quickly to changes
- Brings nice properties

Part II: safe shallow model-based RL

Direct RL



RL towards control theory - the BO of RL